

Deep Reinforcement Learning for Modeling User Strategies and Motivations from Massive Game Behavioral Data

Leave Authors Anonymous

for Submission
City, Country
e-mail address

Leave Authors Anonymous

for Submission
City, Country
e-mail address

Leave Authors Anonymous

for Submission
City, Country
e-mail address

Leave Authors Anonymous

for Submission
City, Country
e-mail address

ABSTRACT

Game user research, which investigates interactions between players and games, is one of the most popular topics in HCI field. However, previous studies are mostly exploratory and descriptive, failing to capture the richness and complexity of game user behaviors. Inspired by recent success of reinforcement learning (RL) algorithm in game strategies, we present a novel RL method to model the underlying motivations which induces the user behaviors we observed. It further reveals the playing strategies which is used to predict user behaviors. We use a massive and complex online game behavioral dataset, World of Warcraft Avatar History (WoWAH), which records over 70,000 users' log data spanning 3 years time period. Our model not only can predict users' behaviors with a high level of accuracy. Moreover, it can also reveal users' motivation dynamics, in terms of achievement, social, and immersion.

ACM Classification Keywords

H.5.m Information Interfaces and Presentation (e.g. HCI): Miscellaneous; I.2.1 Artificial Intelligence: Applications and Expert Systems

Author Keywords

Online game motivation; game design; reinforcement learning; apprenticeship learning;

INTRODUCTION

Understanding Massively multiplayer online role-playing game (MMORPG) games satisfaction mechanism and user behaviors could be non-trivial. As human players

Paste the appropriate copyright statement here. ACM now supports three different copyright statements:

- ACM copyright: ACM holds the copyright on the work. This is the historical approach.
- License: The author(s) retain copyright, but ACM receives an exclusive publication license.
- Open Access: The author(s) wish to pay for the work to be open access. The additional fee must be paid to ACM.

This text field is large enough to hold the appropriate release statement assuming it is single spaced.

Every submission will be assigned their own unique DOI string to be included here.

Attribute	Value
Duration	1107 days
Sample Interval	Every 10 minutes
#Users	70,000+
Locations	One of 165 zones

Table 1. WoWAH Dataset Attributes

have a mixed feeling from different perceptions and they act not for a concrete, explicit objective such as winning an episode or taking high scores. While it's plausible to build a gaming bot from the player log data, to advance in the game, it ignores other dimensions of motivation which the players also care about. The game designers and researchers, also parents and educators, are, however, keen to reveal the underlying mechanism instead of just mastering the game.

Take World of Warcraft (WoW), which is one of the most successful MMO games in the world, as an example. It's massive, and multiplayer as millions of players pay to subscribe the game platform. On the platform players can communicate with others, finish some quests cooperatively, compete with each other, and even build their own guilds. WoW also provides different kinds of races, careers and optional requests that could serve players extra flexibility to evolve different game styles. With all those features the players' behavior can be very complex, as they receive different dimensions of satisfaction which eventually compose the general reward system of the game.

For its huge research and commercial value to study game players' behavior, attentions has been drawn for both qualitative and quantitative studies. In contrary to WoW, many studies are conducted on single-player game or multi-player but zero-sum game. For those environments, advancing usually serves as the sole criterion of game motivation. In Game Trace Archive (GTA) [6] who has been collecting online gaming data since 2012, 10 out of 14 traces are from the game environment

stated above. Although reinforcement learning (RL) algorithms have been succeeded in mastering such kind of games, it does not help figuring out the underlying motivation, neither the complex reward mechanism existing in MMO games and in real world. Even for those studies in WoW, datasets such as Very Large WoW Armory Dataset [4] focus on certain campaigns instead of general gaming strategies because of its short time span and fine-grained actions (such as spelling). Taking that into consideration, the episodes recorded is still mostly about advancing, and game tactics.

World of Warcraft Avatar History (WoWAH) dataset [8], published in 2011, records the location (in term of zone) of over 70,000 users every 10 minutes, together with some essential statistical analysis. However, those analysis and the further studies based on the dataset are far from fully utilization of the data. While some of them use simple classifiers or clustering for behavior analysis [17, 5], others focus on forecasting future events such as unsubscribing of game or violation of terms [3, 18, 9]. Few of them are effectively modeling the interaction between the players and the game environment, neither on another dataset. On the other hand some studies tries to conduct studies on player motivations [4], but only able to apply some basic machine learning tools such as clustering.

We re-employ the WoWAH dataset in a totally different way: to model human-computer interaction (HCI) from a reinforcement learning perspective. The human player and computer are characterized by a pair of agent and environment, as shown in Fig. 1. We show an interesting analogy between the traditional user experience model on the left, and the reinforcement learning scheme on the right. In the model, each player is modeled by an agent, whose zone transactions are regarded as the actions of the agent. Instead of manually collect opinions from the players as the feedback, we utilize the actual reward system of the game which provide diverse satisfactions. A common difficulty in WoW and most of the cases in HCI is that the definition of reward is implicit or partial, e.g. we can't tell if a play is optimal simply by looking at the experience gain credited to this play. Even if some could try building a reward function according to their knowledge of certain environments, it worth to note that figuring out the reward mechanism itself is a valuable task. In fact, the, usually underlying, reward mechanism gives the most succinct description of the task. Especially for the developers of the environment, it's a way to give a quantitative representation of what they have designed.

The process of recovering the reward function is inverse reinforcement learning (IRL). In a data-driven manner, it recovers the underlying reward function which induces the recorded user behaviors, simply by assuming that those players are trying to maximize their rewards. Most of the IRL algorithms require we have access the dynamic of the environment, that is in WoW, we simulate

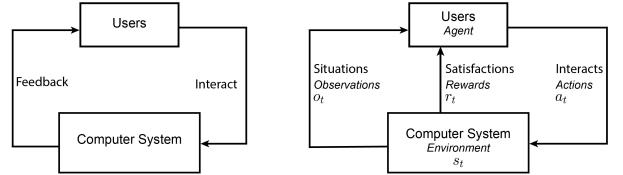


Figure 1. User experience model (left) and reinforcement learning model (right)

the player's action and observe the feedback from the environment. Although we obviously don't have it, we address it by proposing an IRL algorithm based on purely off-policy learning. The model analyzes the behaviors of the players, using the trajectories composed of player locations, and returns the reward function under which those trajectories are most plausible to happen. With the recovered reward, we train an agent to best mimic the general user behavior in WoWAH dataset. As our reward recovery algorithm works for any set of trajectories, we conduct future studies to character the dynamics of the game environment thus how design factors working on user behaviors, and compare different groups of users to examine the divergence of game motivations. Consistency was observed between our model results and empirical analysis on the game environment.

RELATED WORK

Online Game Player Motivation And Rewards

One reason that online games appeal so many players is that it provides the possibility for different kinds of play styles. The same online game may have respective meanings for different players. In the game environment, players pursue certain kinds of satisfactions they want by conducting their actions, which eventually compose the user behaviors we observe. Hence, it's natural to assume a user's behaviors are highly correlated with the user's specific reward mechanism, represented by their demanding satisfactions. And for game designers and research, finding the demanding satisfactions of users could yield better understanding of the game, and ease to serve the users with better experience.

Previous study categories the satisfactions into three major aspects, namely, achievement satisfactions, social satisfactions, and immersion satisfactions [20]. It was then divided into ten subcategories, briefed in Table 2. More details are provided in the original paper and the author's another paper about MMORPG motivation [19]. The satisfaction model is popular in online game researches, and provide some basic understand of the online game motivation. However, the weights associated with these ten kinds of satisfactions are still unknown. As the weights basically depicts a user's profile, it would be very meaningful if they can be recovered in a purely data-driven process, i.e. based on the user's playing history.

Components	Sub-components
Achievement	Advancement, Mechanics, Competition
Social Immersion	Socializing, Relationship, Teamwork Discovery, Role Playing, Customization, Escapism

Table 2. Components of game motivation

Player Behavior Analysis in Online Games

Reinforcement Learning and Inverse Reinforcement learning

Reinforcement learning (RL), especially deep reinforcement learning (DRL), is an emerging domain inspired by behaviorist psychology. In RL, the agent performs in an online environment and conducts a sequence of actions. Instead of being taught of the correct policy, the agent gets a reward for each of its actions and evolves itself from its own pains and pleasures. Interestingly, it deals with the common trade-off between exploration and exploitation, which presents in many real life situations. Success in maximizing cumulative reward in the future, rather than exploiting current reward, could lead to expert level game play, which is also a intrinsic property of expert human players. For example, to level up quickly in a game in a long-term perspective, the player has to conduct some preparation works which does not necessarily benefit its leveling up in the near future.

RL is able to formulate many problems in user behaviors analysis and modeling. For example, the recent advances in RL conduct superior play compared with professional human players in the game of Go [16], Atari 2600 [11], and Poker [7]. It's also employed to model users' clicking and browsing behaviors for online shopping websites, and robotics and optimal controls. It's worth to note that DRL makes it possible to handle complex environment with high-dimension observation spaces, making it possible to model a wide range of problems such as WoW user behavior analysis.

A great challenge is that the reward scheme of the environment is not explicitly given in most problems in user behavior analysis and modeling. For example, the user experience of a software or an online game is the combination of different kinds of satisfaction, processed by the human brains. In some of those cases, inverse reinforcement learning (IRL) [13, 12] recovers the underlying reward function, using the recorded actions of the users. Further applying apprenticeship learning algorithms [1] could mimic the policy that was used to generate the recorded data. But to invoke those IRL algorithms, it usually requires the access of the dynamic of the environment, which is not doable in online gaming analysis. As applying IRL on users' behavior log could help the developers understand the intention of the user, New IRL algorithms are needed to be designed to comply with the problem without the access to the dynamics.

METHOD

Dataset Description and Reinforcement Learning Formulation

We use World of Warcraft Avatar History (WoWAH) dataset [8], collected from realm *TW-Light's Hope* during 1st Jan 2006 to 10th Jan 2009. It contains 70,055 users after we filtering out those with too short playing history, each with 516.21 time intervals (5162.1 minutes) spent online on average. During each of the 10-minute interval, the user's location, in terms of zone ID (ranging from 1 to 165), was recorded. The players' trajectories was composed of a sequence of locations, which reflect their playing strategies. We find datasets with coarse-grain records, i.e. 10-minute interval, most suitable to model user behaviors. They are preferred than those with interval of a few seconds, which instead focus on game tactics and user mobility [4, 15]. Moreover, the dataset contains many different kinds of users: both novice and expert, guild members and isolated players, low-level and high-level players etc. with their respect class, and race in game. The detailed attributes are listed in Table reftbl:-wowah.

In a reinforcement learning perspective, we treat an user as the agent who conduct an action every time interval. An action here corresponds to a zone transaction or staying in the current zone. The user receives current stats of its profile as the observation, after each action it has conducted. The observation includes all the attributes recorded in the dataset, cached by a certain number of time intervals, after a pre-processing to reduce the redundancy. To make the data consistent, we ignore the time intervals where a user being offline, and concatenate multiple game session of a player to build the trajectory. As the exact reward mechanism remains unknown, the RL policy can only be learned after the reward function is properly modeled.

We process some preparation for the reward modeling. Borrowing the result in game motivation analysis [20], we define and compute different of satisfactions using the WoWAH dataset. The constructions of f_1, \dots, f_5 are illustrated in Table 3, which are able to be calculated from the observations the user receives. With those value, we assume the reward function which the agent tries to optimize during the game play, is associated with the five different dimensions of satisfactions. We further address the relation between the reward function, and the satisfactions. Note the value of f_1, \dots, f_5 are normalized into the same scale.

Reward Mechanism Modeling

We present our algorithm to recover the underlying reward mechanism of the players, using the trajectory τ of a user or a group of users. Assume the reward a user receives is the convex combination of the five kinds of satisfactions listed in Table. 3. Let $f_t = (f_t^1, f_t^2, f_t^3, f_t^4, f_t^5)$ be the vector of satisfactions the user received at time t ,

Sat. Category & Definition

f^1	Advancement The speed the user collecting experience and leveling up in game. It's the most common satisfaction a user could receive.
f^2	Competition The satisfaction the user get by joining battleground or arena and competing with human opponents.
f^3	Relationship The long-term relationship with the user's guild, which is quantified by the time elapse since the user joining its guild.
f^4	Teamwork The satisfaction the user get by playing in a zone which is featured by teamwork, e.g. <i>Battleground</i> , <i>Arena</i> , <i>Dungeon</i> , <i>Raid</i> , or a zone controlled by <i>The Alliance</i> .
f^5	Escapism Escapism begins to accumulate if the user has been online for a long session or has been regularly login to the game for many days.

Table 3. Different types of satisfactions (Sat.), and their corresponding categories

we have the total reward the user receives at time t

$$r_t = \phi^T f_t,$$

where ϕ is the combination weight with $\|\phi\|_1 = 1$, $\phi \geq 0$. Assume at each time step, the user tries to take an action a according to the current game state s so as to maximize the best expected cumulative discounted (with discount γ over time) reward $Q^*(s, a)$, where

$$Q^*(s, a) = \mathbf{E}[R_t | s_t = s, a_t = a | \pi^*]$$

and

$$R_t = \sum_{t' \geq t} \gamma^{t'-t} r_{t'} = \sum_{t' \geq t} \gamma^{t'-t} \phi^T f_{t'}.$$

The term π^* indicates optimal policy, described by a probabilistic distribution $\mathbb{P}(a|s)$ over the feasible action space $\mathcal{A}(s)$, and Q^* is also known as the action-value function. In this setting, the weight ϕ must satisfy that the action the user has taken must induce a larger Q^* value than any other valid action would have done. This optimality infers that

$$Q^*(s, a) \geq \max_{a' \in \mathcal{A}(s)} Q^*(s, a') \quad (1)$$

is satisfied for all (s, a) pairs appeared in the user's trajectory. Consider the existence of possible sub-optimal actions, we introduce slack variables $\xi_{s,a}$ into the Eq. (1) to make it feasible. Let $\xi_{s,a}$ be the difference of the actual action-value $Q^*(s, a)$, and the largest possible action-value $\max_{a' \in \mathcal{A}(s)} Q^*(s, a')$ whenever Eq. (1) is not satisfied, and zero otherwise. We minimize the summation of $\xi_{s,a}$ over the recorded user's trajectory

$$-C \sum_{s,a} \left[\min(0, Q^*(s, a) - \max_{a' \in \mathcal{A}(s) \setminus a} Q^*(s, a')) \right], \quad (2)$$

which is then formulated into the following linear programming (LP) problem

$$\begin{aligned} & \underset{\phi, \xi}{\text{minimize}} \quad C \sum \xi_{s,a} \\ & \text{subjectto} \quad \phi^T (\tilde{Q}(s, a) - \tilde{Q}(s, a')) \geq -\xi_{s,a}, \\ & \quad \forall(s, a) \in \tau, a' \in \mathcal{A}(s) \setminus a \\ & \quad \phi \geq 0, \|\phi\|_1 \geq 1 \\ & \quad \xi_{s,a} \geq 0, \forall(s, a) \in \tau. \end{aligned} \quad (3)$$

In formula (3), $\tilde{Q}(s, a) = (Q^1, Q^2, Q^3, Q^4, Q^5)^T$, and

$$Q^i(s, a) = \mathbf{E}[\sum_{t' \geq t} \gamma^{t'-t} f_{t'}^i | s_t = s, a_t = a | \pi^{i,*}] \quad (4)$$

is the action-value function, when the user only takes the $r = f^i$ into account and ignores all 4 other kinds of satisfactions. The LP formulated above is equivalent to minimizing Eq. (2) because by definition we have $Q^*(s, a) = \phi^T \tilde{Q}(s, a)$

Action-Value Function Approximation

As LP can be solved efficiently, it suffices to estimate $Q^i(s, a)$. As the number of possible states s could be arbitrary large for an user in WoW, it's impossible to enumerate over the state space as typical Markov Decision Process (MDP) does. Instead, we build deep-Q networks (DQN) [11] to approximate the Q^i functions. The network takes the observation as input which is used to represent the state s , and output the Q value for every action a . Taking advantage of the fact that DQN is a general framework for multiple environments, we use the same network architecture for $i = 1, \dots, 5$ as illustrated in Fig. 3. The categorical elements in s are firstly processed by an embedding layer [10], while the numeral elements are fed into an fully connected (FC) layer with rectifier non-linearity. The output of embedding layer and FC layer are then concatenated and fed into another FC layer with rectifier non-linearity. A final FC layer is applied to compute the $Q(s, a)$ value for each action $a \in \cup_s \mathcal{A}(s)$.

Denote the trainable parameters in the Q-network as θ^i , we optimize over θ^i in order to approximate Eq. (4). The key observation of Q-learning is that, the action-value function, by its definition, should satisfy the Bellman equation. That is, if the user takes action a and the state turns into s_{t+1} from s_t , we have (recall f_t^i is the reward associated with Q^i)

$$Q^i(s_t, a_t) = f_t^i + \gamma \max_{a'} Q^i(s_{t+1}, a'). \quad (5)$$

Q-learning tries to find the action-value function satisfying Eq. (5), by minimizing the squared difference L between both sides of the Bellman equation

$$L^i = \mathbb{E}_{s_t, a_t, f_t^i, s_{t+1}} \frac{1}{2} (Q^i(s_t, a_t) - f_t^i - \gamma \max_{a'} Q^i(s_{t+1}, a'))^2.$$

Since L^i is differentiable with respect to θ^i , θ^i can be updated via stochastic gradient descent, by

$$\theta^i \leftarrow \theta^i - \alpha \frac{\partial L}{\partial \theta^i} \Big|_{s_t, a_t, f_t^i, s_{t+1}}$$

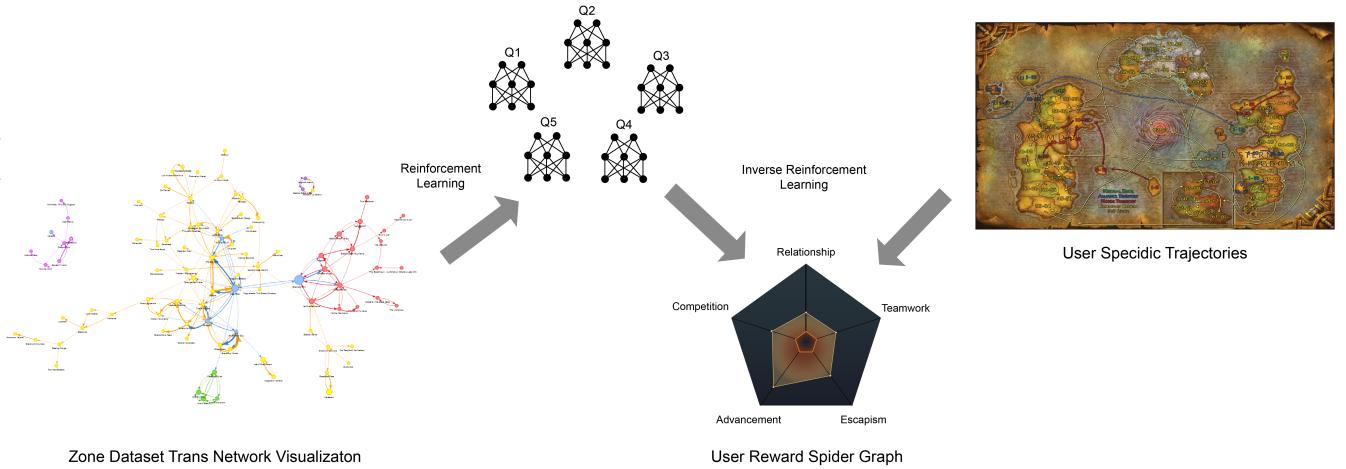


Figure 2. Workflow to recovery the underlying reward mechanism using RL and IRL.

Different from the algorithm introduced in [11], that the trajectories, we try to find if a matches $\pi^1(s)$. The pre-target network $Q(s_{t+1}, a')$ only get updated periodically, diction accuracy is 0.49 with the total number of feasible we put vanilla different of L^i . This is because when using actions $|\cup_s \mathcal{A}(s)| = 165$, corresponding to 165 different log data, we can easily sample from the dataset randomly zones in WoW existed during Jan 2006 - Jan 2009. It's and the instability of DQN caused by the correlation be-competitive with 0.53, using policy cloning [2, 14] with tween consecutive frames can be avoided. The derivative classifier, and outperforms 0.23 if we approximate the Q function using linear function. We conclude the quality of Q^1 is decent.

$$\frac{\partial L}{\partial \theta^i} = \mathbb{E} \left[(Q^i(s_t, a_t) - f_t^i - \gamma \max_{a'} Q(s_{t+1}, a')) \cdot \left(\frac{\partial Q^i(s, a)}{\partial \theta^i} |_{s_t, a_t} - \frac{\partial Q^i(s, a)}{\partial \theta^i} |_{s_{t+1}, a^*} \right) \right],$$

where a^* is the action which maximizes $Q^i(s_{t+1}, a)$.

Policy and User Behaviors

Solving LP (3), the action-value function $Q^*(s, a)$ overunderlying reward mechanism of the player community the underlying mechanism is known for the user/usersis revealed as $\phi = (0.61, 0.16, 0.07, 0.05, 0.10)^T$. In an who conducted trajectories τ . As assumed by the def-other words, when conducting a behavior, in general the inition of $Q^*(s, a)$, player pursuing cumulative rewarduser's consideration is composed of 61% there advancement- R_t should follow the action a' which leads to largestement, 16% their competition, 7% their relationship, 5% their teamwork, and 10% their escapism. The results is $Q^*(s, a')$. Which induces the policy illustrated as a spider map in Fig. ??.

$$\pi^*(s) = \operatorname{argmax}_{a'} Q^*(s, a').$$

Although $\pi^*(s)$ may not be the one who advances fastest in game (which should instead be $\pi^1(s) = \operatorname{argmax}_{a'} Q^1(s, a')$), it is able to mimic the actions that in general a player would conduct. That is, the policy $\pi^*(s)$ can be used to predict the action of users, even when they are under mixed and complex reward mechanism.

Using similar idea, during the training of Q^1 network,derlying reward machenism between different groups of we evaluate the quality of the Q-learning process. Con-people, which is the causality driving the divergence of sider collecting experience and getting level up is one of user behaviors. For example as the level of a player be- the major objective for players, we evaluate if the ac-comes higher, it tends to cares more about relationship tions predicted by Q^1 agree with the moves conductedthan advancement. Finally we show, using our model, by those who level up quickly. For the (s, a) pairs ex-the evolution of underlying user satisfaction over time. tracted from the top 200 (in leveling up speed) players' Especially, observing the dynamics around the release

RESULTS

Different from previous approaches on Atari games [11] and zero-sum games [16, 7], our approach models the user behavior instead of simply pursuing advancement. Using Table. 3 and solve LP (3) on a dataset randomly drawn from the whole WoWAH dataset, the universal underlying reward mechanism of the player community is known for the user/usersis revealed as $\phi = (0.61, 0.16, 0.07, 0.05, 0.10)^T$. In an who conducted trajectories τ . As assumed by the def-other words, when conducting a behavior, in general the inition of $Q^*(s, a)$, player pursuing cumulative rewarduser's consideration is composed of 61% there advancement- R_t should follow the action a' which leads to largestement, 16% their competition, 7% their relationship, 5% their teamwork, and 10% their escapism. The results is illustrated as a spider map in Fig. ??.

Armed with the model we conduct three different kinds further studies on WoWAH dataset. First we show that our model better describes the user behavior, than those who takes only single satisfaction metric. After recovering the underlying reward mechanism of the environment, our model could predict the move of the users, instead of just suggesting the best move for advancing (as Q^1 does). We then show the divergence of the un-

derlying reward mechanism between different groups of we evaluate the quality of the Q-learning process. Con-people, which is the causality driving the divergence of

sider collecting experience and getting level up is one of user behaviors. For example as the level of a player be- the major objective for players, we evaluate if the ac-comes higher, it tends to cares more about relationship tions predicted by Q^1 agree with the moves conductedthan advancement. Finally we show, using our model, by those who level up quickly. For the (s, a) pairs ex-the evolution of underlying user satisfaction over time. tracted from the top 200 (in leveling up speed) players' Especially, observing the dynamics around the release

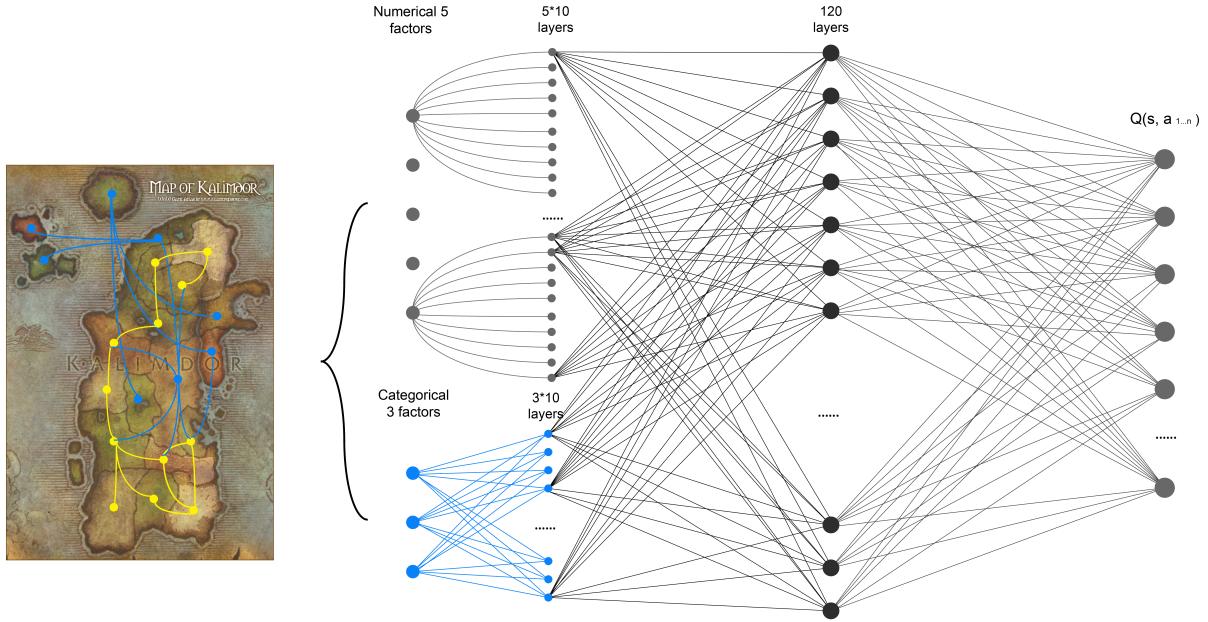


Figure 3. DQN architecture for Q^i training, $i = 1, \dots, 5$




Figure 4. Workflow to generate our results

of patch *Fall of the Lich King*¹, it shed some light on quantitative models of the game design.

Predicting the Users' Behavior

explain different game play styles

Armed with the reward mechanism we are able to model users' behavior instead of creating an agent to play the game itself. Consider that the users' motivation of playing is more than leveling up,

Another interesting observation is the tradeoff between exploitation and exploration. In fact, many users' actions are sub-optimal, in the sense that they care too much about the satisfaction in the near future.

¹http://wowwiki.wikia.com/wiki/Patch_3.3.0

Behavior Evolution over Time

The way the user

Behavior Divergence between Different Groups of User

REFERENCES

1. Pieter Abbeel and Andrew Y Ng. 2004. Apprenticeship learning via inverse reinforcement learning. In *Proceedings of the twenty-first international conference on Machine learning*. ACM, 1.
2. R Amit and Maja J Mataric. 2002. Parametric primitives for motor representation and control. In *Robotics and Automation, 2002. Proceedings. ICRA'02. IEEE International Conference on*, Vol. 1. IEEE, 863–868.
3. Christian Bauckhage, Anders Drachen, and Rafet Sifa. 2015. Clustering game behavior data. *IEEE Transactions on Computational Intelligence and AI in Games* 7, 3 (2015), 266–278.
4. Jonathan Bell, Swapneel Sheth, and Gail Kaiser. 2013. A large-scale, longitudinal study of user profiles in world of warcraft. In *Proceedings of the 22nd international conference on World Wide Web companion*. International World Wide Web Conferences Steering Committee, 1175–1184.
5. Anders Drachen, Christian Thurau, Rafet Sifa, and Christian Bauckhage. 2014. A comparison of methods for player clustering via behavioral telemetry. *arXiv preprint arXiv:1407.3950* (2014).

6. Yong Guo and Alexandru Iosup. 2012. The game trace archive. In *Proceedings of the 11th Annual Workshop on Network and Systems Support for Games*. IEEE Press, 4.
7. Johannes Heinrich and David Silver. 2016. Deep Reinforcement Learning from Self-Play in Imperfect-Information Games. *arXiv preprint arXiv:1603.01121* (2016).
8. Yeng-Ting Lee, Kuan-Ta Chen, Yun-Maw Cheng, and Chin-Laung Lei. 2011. World of Warcraft avatar history dataset. In *Proceedings of the second annual ACM conference on Multimedia systems*. ACM, 123–128.
9. Jing-Kai Lou, Kuan-Ta Chen, Hwai-Jung Hsu, and Chin-Laung Lei. 2012. Forecasting online game addictiveness. In *Proceedings of the 11th Annual Workshop on Network and Systems Support for Games*. IEEE Press, 6.
10. T Mikolov and J Dean. 2013. Distributed representations of words and phrases and their compositionality. *Advances in neural information processing systems* (2013).
11. Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex Graves, Martin Riedmiller, Andreas K Fidjeland, Georg Ostrovski, and others. 2015. Human-level control through deep reinforcement learning. *Nature* 518, 7540 (2015), 529–533.
12. Andrew Y Ng, Stuart J Russell, and others. 2000. Algorithms for inverse reinforcement learning.. In *Icml*. 663–670.
13. Nathan D Ratliff, J Andrew Bagnell, and Martin A Zinkevich. 2006. Maximum margin planning. In *Proceedings of the 23rd international conference on Machine learning*. ACM, 729–736.
14. Claude Sammut, Scott Hurst, Dana Kedzier, Donald Michie, and others. 1992. Learning to fly. In *Proceedings of the ninth international workshop on Machine learning*. 385–393.
15. Siqi Shen, Niels Brouwers, Alexandru Iosup, and Dick Epema. 2014. Characterization of human mobility in networked virtual environments. In *Proceedings of Network and Operating System Support on Digital Audio and Video Workshop*. ACM, 13.
16. David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, and others. 2016. Mastering the game of Go with deep neural networks and tree search. *Nature* 529, 7587 (2016), 484–489.
17. Mirko Suznjevic, Ivana Stupar, and Maja Matijasevic. 2011. MMORPG player behavior model based on player action categories. In *Proceedings of the 10th Annual Workshop on Network and Systems Support for Games*. IEEE Press, 6.
18. Ruck Thawonmas, Keisuke Yoshida, Jing-Kai Lou, and Kuan-Ta Chen. 2011. Analysis of revisitations in online games. *Entertainment Computing* 2, 4 (2011), 215–221.
19. Nick Yee. 2006a. The demographics, motivations, and derived experiences of users of massively multi-user online graphical environments. *Presence: Teleoperators and virtual environments* 15, 3 (2006), 309–329.
20. Nick Yee. 2006b. Motivations for play in online games. *CyberPsychology & behavior* 9, 6 (2006), 772–775.